

Lakelet

Executive summary · Pre-seed · September 2026

Your data fits on your laptop.

Open-source, local-first lakehouse that anyone can ask questions of. Borrows a bigger computer only for the runs that need one.

\$15B

cloud data warehouse spend, 2026

~\$3B

small teams under a terabyte (est.)

\$1.5M

pre-seed raise, ~18 months

\$0

to start; free local tier forever

The problem

Most teams pay by the second for a computer they already own. A five-person team with 38 GB of data runs a Snowflake X-Small nine hours a day and pays ~\$4,200 a month for queries a MacBook would finish in seconds. Snowflake enterprise averages \$745K a year; even SMB tiers run \$59–185K. Every incumbent raised its floor in the last 18 months: Databricks retired its Standard tier, MotherDuck removed its \$25 plan.

Why now

Three open-source tools matured within months of each other in 2026. Apache Polaris graduated and DuckLake 1.0 shipped (Feb–Apr), removing the catalog problem. DuckDB 1.5.3 shipped Iceberg writes (May 29), removing the Spark cluster that 96% of Iceberg users ran just to write a table. AWS acquired the DuckLabs team (Aug 26) while the project stays MIT under an independent foundation. Meanwhile the desk got bigger: a MacBook Pro M5 Max holds 128 GB in unified memory at 614 GB/s, and a \$1,699 Mac mini carries 64 GB.

The product

One box. Ask in English or write SQL; either way the SQL is shown and editable, and nothing runs before the gauge. **The gauge.** A pre-flight verdict on every query: Green runs here, Yellow runs slowly, Red needs more machine. **Burst.** One click sends a Red run to an ephemeral worker in your own cloud account with a hard cost cap the worker enforces. **Simple / Technical.** Every project is a dbt project under git from the first save; non-developers see questions, checks and versions; engineers see the DAG, the SQL and the diff. **Agents.** An MCP server exposes the same tools to Claude Code, Cursor and Codex, with the gauge as a guardrail and a per-agent daily cap the worker enforces.

How it runs

Data files stay in the customer's S3, GCS or R2 bucket in every setup. Lakelet's cloud holds catalog metadata only. Solo users run everything on the laptop with an SQLite catalog; teams share a hosted Iceberg REST catalog on Postgres; companies deploy the same catalog and control plane into their own VPC.

Business model

The Supabase playbook: an open core on a standard, free to start, priced like a tool, provisioned by AI tools on the user's behalf. Over 60% of new Supabase databases are now launched by an AI tool; Lakelet ships an MCP server on day one and a platforms API on the Team tier so app builders can give their users a lakehouse without hosting one. Local app and CLI are Apache 2.0 and free forever. Team catalog at \$24 per seat per month. Burst at cost plus 15% with a per-run cap. Self-hosted from \$25K a year from 2029. The defensible asset is the calibration data every gauge run produces; the catalog is where switching cost lives.

Competition

MotherDuck (~\$133M raised, ~\$496M valuation) is DuckDB on their cloud; Lakelet is DuckDB on your bucket, the Supabase-versus-Firebase distinction. Snowflake (\$4.7B FY26, ~11,000 customers, 733 over \$1M) and Databricks (\$4.8B run rate) win on elastic scale and lose on cost, locality and lock-in below a terabyte. AWS may ship a serverless DuckDB; Lakelet's answer is multi-cloud, an open catalog and the UX, not the engine.

Plan and milestones

Day 0 · Sep–Dec 2026	Desktop app, CLI, gauge, burst on AWS, MCP server. Ten design partners on real workloads. Pre-seed raised on this.
Day 1 · H1 2027	Open-source launch. 2,000 monthly actives; 30 teams waitlisted with a real bill.
Day 2 · to Jun 2028	Team catalog, billing, platforms API. ~200 paying teams, ~\$300K ARR. Seed (~\$4M).
Day 3 · to Dec 2029	Self-hosted, SSO, SOC 2, connectors. \$1–3M ARR. Series A.

Financials (base case, end-of-year ARR)

2027: \$49K · 2028: \$326K · 2029: \$1.35M, on 8K → 40K → 120K monthly actives and 1.5–2.5% seat conversion (open-source benchmark 1–3%). Bull case \$7M by 2029 if coding agents or an AI app builder default to provisioning Lakelet, as they did for Supabase. Cumulative burn through 2029 about \$3M.

Exit paths

Strategic buyers in this layer pay for position in an open standard and for the team, not revenue: Tabular (>\$1B, ~\$1M ARR), Neon (~\$1B), Crunchy Data (~\$250M on \$30M+ ARR), DuckLabs (undisclosed, no VC). Likeliest acquirers: Fivetran + dbt Labs, Snowflake, Databricks, Cloudflare, Google. Independent path: the Supabase curve (\$10.5B).

The ask

\$1.5M pre-seed on a SAFE to launch Lakelet open source and reach the team product in 18 months.

Use of funds: two engineers now, a third plus developer relations in 2027, cloud credits, the open-source launch. No sales hire before 2029. Milestones: ten design partners and a gauge within 2x on 80% of runs by Dec 2026; 2,000 monthly actives by Jun 2027; ~200 paying teams and a seed by mid-2028.

hello@lakelet.dev

Deck, financial plan, architecture and product specs available on request.